

Credit Card Fraud Detection

Puneeth Kondarasi, Gagan Karnam

2025-05-13

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(ranger) #random forest
```

```
## Warning: package 'ranger' was built under R version 4.3.3
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.3.3
```

```
## Loading required package: ggplot2
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
## Loading required package: lattice
```

```
library(caTools)
```

```
## Warning: package 'caTools' was built under R version 4.3.3
```

```
library(data.table)
```

```
## Warning: package 'data.table' was built under R version 4.3.3
```

```
##  
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':  
##  
##   between, first, last
```

```
library(ggplot2)  
library(ROSE) #Data balancing
```

```
## Warning: package 'ROSE' was built under R version 4.3.3
```

```
## Loaded ROSE 0.0-4
```

```
library(pROC) #ROC
```

```
## Warning: package 'pROC' was built under R version 4.3.3
```

```
## Type 'citation("pROC")' for a citation.
```

```
##  
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':  
##  
##   cov, smooth, var
```

```
library(rpart) #decision tree
```

```
## Warning: package 'rpart' was built under R version 4.3.3
```

```
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 4.3.3
```

```
library(xgboost)
```

```
## Warning: package 'xgboost' was built under R version 4.3.3
```

```
##  
## Attaching package: 'xgboost'
```

```
## The following object is masked from 'package:dplyr':  
##  
##   slice
```

```
library(ROCR)
```

```
## Warning: package 'ROCR' was built under R version 4.3.3
```

```
dataset <- setDT(read.csv("C:/Users/punee/Downloads/creditcard.csv"))
```

```
# Data exploration
```

```
head(dataset)
```

```
##      Time      V1      V2      V3      V4      V5      V6
##      <num>      <num>      <num>      <num>      <num>      <num>      <num>
## 1:      0 -1.3598071 -0.07278117 2.5363467 1.3781552 -0.33832077 0.46238778
## 2:      0 1.1918571 0.26615071 0.1664801 0.4481541 0.06001765 -0.08236081
## 3:      1 -1.3583541 -1.34016307 1.7732093 0.3797796 -0.50319813 1.80049938
## 4:      1 -0.9662717 -0.18522601 1.7929933 -0.8632913 -0.01030888 1.24720317
## 5:      2 -1.1582331 0.87773675 1.5487178 0.4030339 -0.40719338 0.09592146
## 6:      2 -0.4259659 0.96052304 1.1411093 -0.1682521 0.42098688 -0.02972755
##      V7      V8      V9      V10      V11      V12
##      <num>      <num>      <num>      <num>      <num>      <num>
## 1: 0.23959855 0.09869790 0.3637870 0.09079417 -0.5515995 -0.61780086
## 2: -0.07880298 0.08510165 -0.2554251 -0.16697441 1.6127267 1.06523531
## 3: 0.79146096 0.24767579 -1.5146543 0.20764287 0.6245015 0.06608369
## 4: 0.23760894 0.37743587 -1.3870241 -0.05495192 -0.2264873 0.17822823
## 5: 0.59294075 -0.27053268 0.8177393 0.75307443 -0.8228429 0.53819555
## 6: 0.47620095 0.26031433 -0.5686714 -0.37140720 1.3412620 0.35989384
##      V13      V14      V15      V16      V17      V18
##      <num>      <num>      <num>      <num>      <num>      <num>
## 1: -0.9913898 -0.3111694 1.4681770 -0.4704005 0.20797124 0.02579058
## 2: 0.4890950 -0.1437723 0.6355581 0.4639170 -0.11480466 -0.18336127
## 3: 0.7172927 -0.1659459 2.3458649 -2.8900832 1.10996938 -0.12135931
## 4: 0.5077569 -0.2879237 -0.6314181 -1.0596472 -0.68409279 1.96577500
## 5: 1.3458516 -1.1196698 0.1751211 -0.4514492 -0.23703324 -0.03819479
## 6: -0.3580907 -0.1371337 0.5176168 0.4017259 -0.05813282 0.06865315
##      V19      V20      V21      V22      V23      V24
##      <num>      <num>      <num>      <num>      <num>      <num>
## 1: 0.40399296 0.25141210 -0.018306778 0.277837576 -0.11047391 0.06692807
## 2: -0.14578304 -0.06908314 -0.225775248 -0.638671953 0.10128802 -0.33984648
## 3: -2.26185710 0.52497973 0.247998153 0.771679402 0.90941226 -0.68928096
## 4: -1.23262197 -0.20803778 -0.108300452 0.005273597 -0.19032052 -1.17557533
## 5: 0.80348692 0.40854236 -0.009430697 0.798278495 -0.13745808 0.14126698
## 6: -0.03319379 0.08496767 -0.208253515 -0.559824796 -0.02639767 -0.37142658
##      V25      V26      V27      V28 Amount Class
##      <num>      <num>      <num>      <num> <num> <int>
## 1: 0.1285394 -0.1891148 0.133558377 -0.02105305 149.62 0
## 2: 0.1671704 0.1258945 -0.008983099 0.01472417 2.69 0
## 3: -0.3276418 -0.1390966 -0.055352794 -0.05975184 378.66 0
## 4: 0.6473760 -0.2219288 0.062722849 0.06145763 123.50 0
## 5: -0.2060096 0.5022922 0.219422230 0.21515315 69.99 0
## 6: -0.2327938 0.1059148 0.253844225 0.08108026 3.67 0
```

```
tail(dataset)
```

##	Time	V1	V2	V3	V4	V5	V6
##	<num>	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	172785	0.1203164	0.93100513	-0.5460121	-0.7450968	1.13031398	-0.2359732
## 2:	172786	-11.8811179	10.07178497	-9.8347835	-2.0666557	-5.36447278	-2.6068373
## 3:	172787	-0.7327887	-0.05508049	2.0350297	-0.7385886	0.86822940	1.0584153
## 4:	172788	1.9195650	-0.30125385	-3.2496398	-0.5578281	2.63051512	3.0312601
## 5:	172788	-0.2404400	0.53048251	0.7025102	0.6897992	-0.37796113	0.6237077
## 6:	172792	-0.5334125	-0.18973334	0.7033374	-0.5062712	-0.01254568	-0.6496167
##	V7	V8	V9	V10	V11	V12	
##	<num>	<num>	<num>	<num>	<num>	<num>	
## 1:	0.8127221	0.1150929	-0.2040635	-0.6574221	0.6448373	0.19091623	
## 2:	-4.9182154	7.3053340	1.9144283	4.3561704	-1.5931053	2.71194079	
## 3:	0.0243297	0.2948687	0.5848000	-0.9759261	-0.1501888	0.91580191	
## 4:	-0.2968265	0.7084172	0.4324540	-0.4847818	0.4116137	0.06311886	
## 5:	-0.6861800	0.6791455	0.3920867	-0.3991257	-1.9338488	-0.96288614	
## 6:	1.5770063	-0.4146504	0.4861795	-0.9154266	-1.0404583	-0.03151305	
##	V13	V14	V15	V16	V17	V18	
##	<num>	<num>	<num>	<num>	<num>	<num>	
## 1:	-0.5463289	-0.73170658	-0.80803553	0.5996281	0.07044075	0.3731103	
## 2:	-0.6892556	4.62694203	-0.92445871	1.1076406	1.99169111	0.5106323	
## 3:	1.2147558	-0.67514296	1.16493091	-0.7117573	-0.02569286	-1.2211789	
## 4:	-0.1836987	-0.51060184	1.32928351	0.1407160	0.31350179	0.3956525	
## 5:	-1.0420817	0.44962444	1.96256312	-0.6085771	0.50992846	1.1139806	
## 6:	-0.1880929	-0.08431647	0.04133346	-0.3026201	-0.66037665	0.1674299	
##	V19	V20	V21	V22	V23	V24	
##	<num>	<num>	<num>	<num>	<num>	<num>	
## 1:	0.1289038	0.0006758329	-0.3142046	-0.8085204	0.05034266	0.102799590	
## 2:	-0.6829197	1.4758291347	0.2134541	0.1118637	1.01447990	-0.509348453	
## 3:	-1.5455561	0.0596158999	0.2142053	0.9243836	0.01246304	-1.016225669	
## 4:	-0.5772518	0.0013959703	0.2320450	0.5782290	-0.03750086	0.640133881	
## 5:	2.8978488	0.1274335158	0.2652449	0.8000487	-0.16329794	0.123205244	
## 6:	-0.2561169	0.3829481049	0.2610573	0.6430784	0.37677701	0.008797379	
##	V25	V26	V27	V28	Amount	Class	
##	<num>	<num>	<num>	<num>	<num>	<int>	
## 1:	-0.4358701	0.1240789	0.217939865	0.06880333	2.69	0	
## 2:	1.4368069	0.2500343	0.943651172	0.82373096	0.77	0	
## 3:	-0.6066240	-0.3952551	0.068472470	-0.05352739	24.79	0	
## 4:	0.2657455	-0.0873706	0.004454772	-0.02656083	67.88	0	
## 5:	-0.5691589	0.5466685	0.108820735	0.10453282	10.00	0	
## 6:	-0.4736487	-0.8182671	-0.002415309	0.01364891	217.00	0	

```
table(dataset$Class)
```

```
##
##      0      1
## 284315   492
```

```
summary(dataset$Amount)
```


```
##      Min.   1st Qu.   Median     Mean  3rd Qu.     Max.
##      0.00     5.60    22.00    88.35    77.17 25691.16
```

```
colSums(is.na(dataset))
```

##	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10
##	0	0	0	0	0	0	0	0	0	0	0
##	V11	V12	V13	V14	V15	V16	V17	V18	V19	V20	V21
##	0	0	0	0	0	0	0	0	0	0	0
##	V22	V23	V24	V25	V26	V27	V28	Amount	Class		
##	0	0	0	0	0	0	0	0	0		


```
hist(dataset$Amount)
```

Histogram of dataset\$Amount



```
# Data visualization
dataset %>%
  ggplot(aes(x = Time, fill = factor(Class))) +
  geom_histogram(bins = 100) +
  labs(x = "Time (seconds)", y = "No. of transactions", title = "Transaction Distribution") +
  facet_grid(Class ~ ., scales = 'free_y') + theme()
```

Transaction Distribution



```
# Feature Scaling
dataset$Amount <- scale(dataset$Amount)

# Prepare dataset
new_data <- dataset[, -c(1)]
new_data$Class <- as.factor(new_data$Class)
levels(new_data$Class) <- c("Not Fraud", "Fraud")

# Train-test split
set.seed(101)
split <- sample.split(new_data$Class, SplitRatio = 0.8)
train_data <- subset(new_data, split == TRUE)
test_data <- subset(new_data, split == FALSE)

# Show train and test data samples
print("Train Data Sample:")
```

```
## [1] "Train Data Sample:"
```

```
head(train_data)
```

##	V1	V2	V3	V4	V5	V6
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	-1.3598071	-0.07278117	2.5363467	1.3781552	-0.33832077	0.46238778
## 2:	1.1918571	0.26615071	0.1664801	0.4481541	0.06001765	-0.08236081
## 3:	-1.3583541	-1.34016307	1.7732093	0.3797796	-0.50319813	1.80049938
## 4:	-0.9662717	-0.18522601	1.7929933	-0.8632913	-0.01030888	1.24720317
## 5:	-1.1582331	0.87773675	1.5487178	0.4030339	-0.40719338	0.09592146
## 6:	-0.4259659	0.96052304	1.1411093	-0.1682521	0.42098688	-0.02972755
##	V7	V8	V9	V10	V11	V12
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	0.23959855	0.09869790	0.3637870	0.09079417	-0.5515995	-0.61780086
## 2:	-0.07880298	0.08510165	-0.2554251	-0.16697441	1.6127267	1.06523531
## 3:	0.79146096	0.24767579	-1.5146543	0.20764287	0.6245015	0.06608369
## 4:	0.23760894	0.37743587	-1.3870241	-0.05495192	-0.2264873	0.17822823
## 5:	0.59294075	-0.27053268	0.8177393	0.75307443	-0.8228429	0.53819555
## 6:	0.47620095	0.26031433	-0.5686714	-0.37140720	1.3412620	0.35989384
##	V13	V14	V15	V16	V17	V18
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	-0.9913898	-0.3111694	1.4681770	-0.4704005	0.20797124	0.02579058
## 2:	0.4890950	-0.1437723	0.6355581	0.4639170	-0.11480466	-0.18336127
## 3:	0.7172927	-0.1659459	2.3458649	-2.8900832	1.10996938	-0.12135931
## 4:	0.5077569	-0.2879237	-0.6314181	-1.0596472	-0.68409279	1.96577500
## 5:	1.3458516	-1.1196698	0.1751211	-0.4514492	-0.23703324	-0.03819479
## 6:	-0.3580907	-0.1371337	0.5176168	0.4017259	-0.05813282	0.06865315
##	V19	V20	V21	V22	V23	V24
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	0.40399296	0.25141210	-0.018306778	0.277837576	-0.11047391	0.06692807
## 2:	-0.14578304	-0.06908314	-0.225775248	-0.638671953	0.10128802	-0.33984648
## 3:	-2.26185710	0.52497973	0.247998153	0.771679402	0.90941226	-0.68928096
## 4:	-1.23262197	-0.20803778	-0.108300452	0.005273597	-0.19032052	-1.17557533
## 5:	0.80348692	0.40854236	-0.009430697	0.798278495	-0.13745808	0.14126698
## 6:	-0.03319379	0.08496767	-0.208253515	-0.559824796	-0.02639767	-0.37142658
##	V25	V26	V27	V28	Amount	Class
##	<num>	<num>	<num>	<num>	<num>	<fctr>
## 1:	0.1285394	-0.1891148	0.133558377	-0.02105305	0.24496383	Not Fraud
## 2:	0.1671704	0.1258945	-0.008983099	0.01472417	-0.34247394	Not Fraud
## 3:	-0.3276418	-0.1390966	-0.055352794	-0.05975184	1.16068389	Not Fraud
## 4:	0.6473760	-0.2219288	0.062722849	0.06145763	0.14053401	Not Fraud
## 5:	-0.2060096	0.5022922	0.219422230	0.21515315	-0.07340321	Not Fraud
## 6:	-0.2327938	0.1059148	0.253844225	0.08108026	-0.33855582	Not Fraud

```
table(train_data$Class)
```

```
##
## Not Fraud      Fraud
##      227452      394
```

```
print("Test Data Sample:")
```

```
## [1] "Test Data Sample:"
```

```
head(test_data)
```

##	V1	V2	V3	V4	V5	V6
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	1.4490438	-1.17633883	0.9138598	-1.3756667	-1.9713832	-0.62915214
## 2:	1.0693736	0.28772213	0.8286127	2.7125204	-0.1783980	0.33754373
## 3:	1.1032154	-0.04029621	1.2673321	1.2890915	-0.7359972	0.28806916
## 4:	0.9624961	0.32846103	-0.1714791	2.1092041	1.1295656	1.69603769
## 5:	-1.9465251	-0.04490051	-0.4055701	-1.0130573	2.9419677	2.95505340
## 6:	-0.5353878	0.86526781	1.3510763	0.1475755	0.4336802	0.08698294
##	V7	V8	V9	V10	V11	V12
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	-1.42323560	0.04845589	-1.7204084	1.6266591	1.19964395	-0.6714398
## 2:	-0.09671686	0.11598174	-0.2210826	0.4602304	-0.77365693	0.3233872
## 3:	-0.58605679	0.18937971	0.7823329	-0.2679751	-0.45031128	0.9367077
## 4:	0.10771161	0.52150216	-1.1913111	0.7243963	1.69032992	0.4067736
## 5:	-0.06306315	0.85554631	0.0499669	0.5737425	-0.08125651	-0.2157450
## 6:	0.69303931	0.17974226	-0.2856419	-0.4824745	0.87179958	0.8534474
##	V13	V14	V15	V16	V17	V18
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	-0.51394715	-0.09504505	0.2309304	0.03196747	0.253414716	0.85434381
## 2:	-0.01107589	-0.17848518	-0.6555643	-0.19992517	0.124005415	-0.98049620
## 3:	0.70838041	-0.46864729	0.3545741	-0.24663466	-0.009212378	-0.59591241
## 4:	-0.93642130	0.98373942	0.7109108	-0.60223177	0.402484376	-1.73716203
## 5:	0.04416063	0.03389776	1.1907177	0.57884348	-0.975667025	0.04406282
## 6:	-0.57182189	0.10225210	-1.5199912	-0.28591250	-0.309633387	-0.40390199
##	V19	V20	V21	V22	V23	V24
##	<num>	<num>	<num>	<num>	<num>	<num>
## 1:	-0.2213654	-0.3872265	-0.009301897	0.3138944	0.02774016	0.5005122871
## 2:	-0.9829161	-0.1531972	-0.036875532	0.0744124	-0.07140743	0.1047437526
## 3:	-0.5756816	-0.1139102	-0.024612006	0.1960020	0.01380165	0.1037583310
## 4:	-2.0276123	-0.2693210	0.143997423	0.4024917	-0.04850822	-1.3718662945
## 5:	0.4886029	-0.2167153	-0.579525934	-0.7992290	0.87030022	0.9834214925
## 6:	-0.8237430	-0.2832638	0.049525687	0.2065365	-0.18710807	0.0007530143
##	V25	V26	V27	V28	Amount	Class
##	<num>	<num>	<num>	<num>	<num>	<fctr>
## 1:	0.25136736	-0.1294780	0.04284987	0.01625326	-0.3220438	Not Fraud
## 2:	0.54826473	0.1040942	0.02149106	0.02129331	-0.2432816	Not Fraud
## 3:	0.36429754	-0.3822606	0.09280919	0.03705052	-0.3012937	Not Fraud
## 4:	0.39081389	0.1999637	0.01637064	-0.01460533	-0.2169343	Not Fraud
## 5:	0.32120113	0.1496499	0.70751884	0.01459975	-0.3496705	Not Fraud
## 6:	0.09811661	-0.5534710	-0.07830550	0.02542738	-0.3461522	Not Fraud

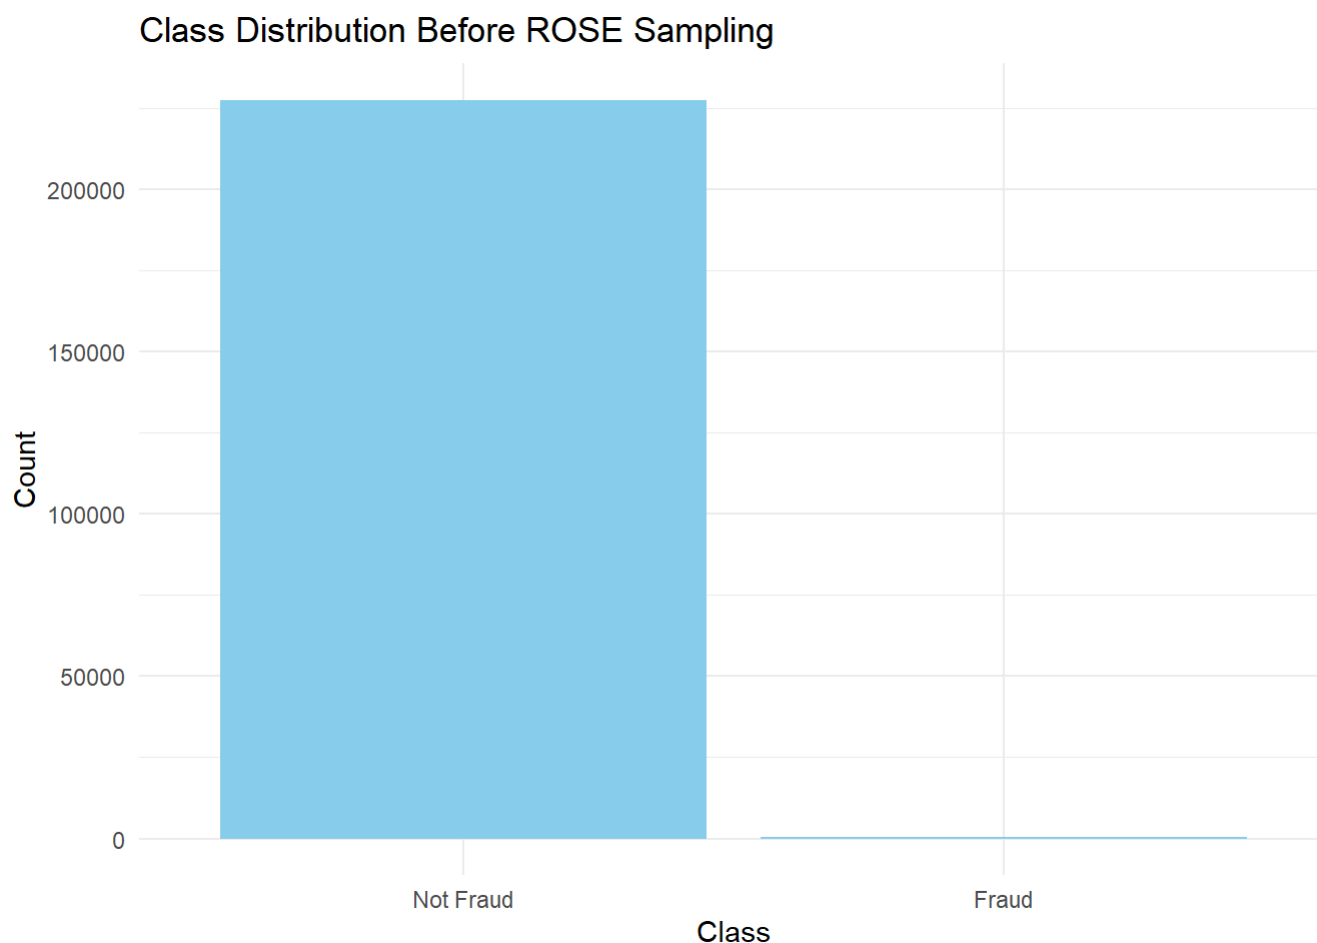
```
table(test_data$Class)
```

```
##
## Not Fraud      Fraud
##      56863      98
```




```
# Handle Class Imbalance using ROSE
set.seed(9560)
rose_train_data <- ROSE(Class ~ ., data = train_data)$data
```

```
# Plot class distribution before ROSE sampling
ggplot(train_data, aes(x = factor(Class))) +
  geom_bar(fill = "skyblue") +
  ggtitle("Class Distribution Before ROSE Sampling") +
  xlab("Class") +
  ylab("Count") +
  theme_minimal()
```



```
# Plot class distribution after ROSE sampling
ggplot(rose_train_data, aes(x = factor(Class))) +
  geom_bar(fill = "lightgreen") +
  ggtitle("Class Distribution After ROSE Sampling") +
  xlab("Class") +
  ylab("Count") +
  theme_minimal()
```

Class Distribution After ROSE Sampling



```
# ----- LOGISTIC REGRESSION -----  
logistic_model <- glm(Class ~ ., data = rose_train_data, family = 'binomial')
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


```
logistic_predictions <- predict(logistic_model, test_data, type = 'response')
```

```
# ROC Curve for Logistic Regression  
roc.curve(test_data$Class, logistic_predictions,  
          plotit = TRUE,  
          col = "blue",  
          main = "ROC Curve - Logistic Regression")
```

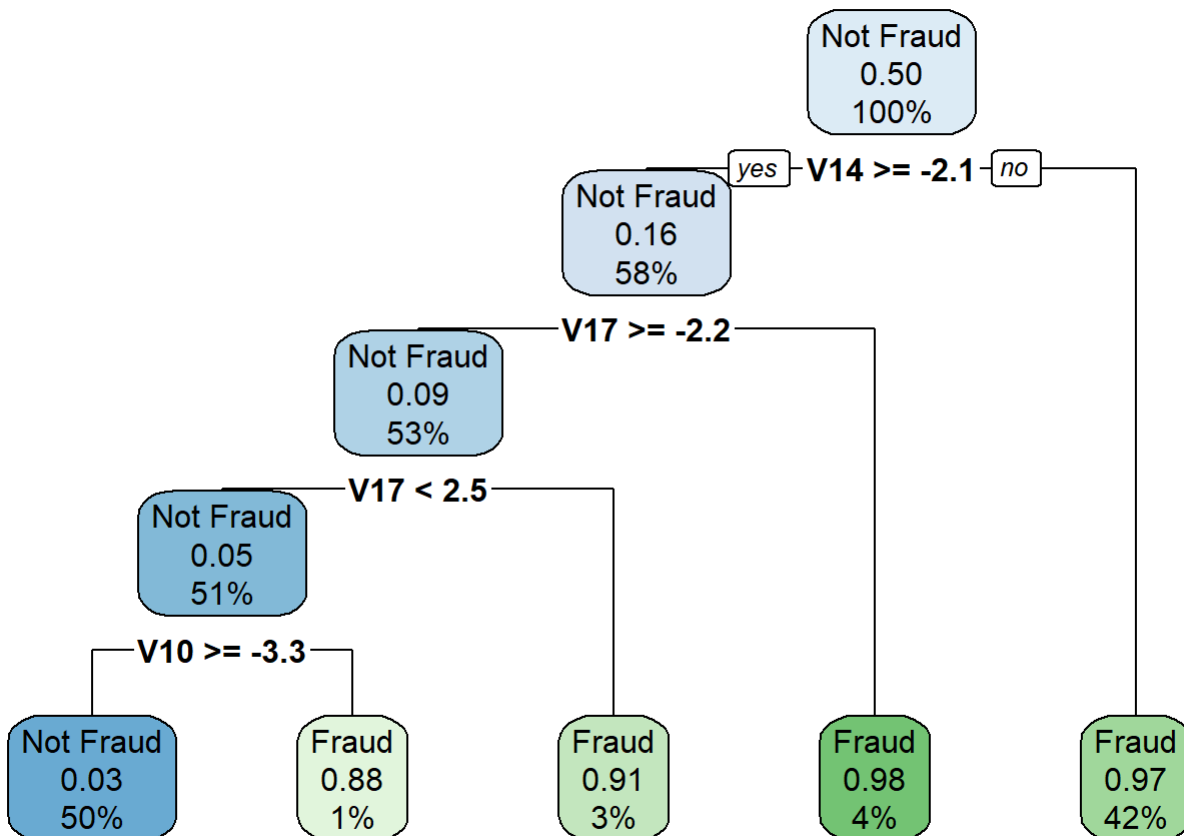
```
## Area under the curve (AUC): 0.977
```

```
legend("bottomright", legend = "Logistic Regression", col = "blue", lwd = 2)
```

ROC Curve - Logistic Regression



```
# ----- DECISION TREE -----  
decisionTree_model <- rpart(Class ~ ., data = rose_train_data, method = 'class')  
rpart.plot(decisionTree_model)
```




```
decisionTree_predictions <- predict(decisionTree_model, test_data, type = 'prob')[,2]
```

```
# ROC Curve for Decision Tree
roc.curve(test_data$Class, decisionTree_predictions,
          plotit = TRUE,
          col = "red",
          main = "ROC Curve - Decision Tree")
```

```
## Area under the curve (AUC): 0.908
```

```
legend("bottomright", legend = "Decision Tree", col = "red", lwd = 2)
```



```
# ----- RANDOM FOREST -----
rf_fit <- ranger(Class ~ .,
                 data = rose_train_data,
                 num.trees = 200,
                 mtry = 5,
                 min.node.size = 10,
                 importance = 'impurity',
                 probability = TRUE)
```


```
## Growing trees.. Progress: 11%. Estimated remaining time: 4 minutes, 41 seconds.
## Growing trees.. Progress: 22%. Estimated remaining time: 4 minutes, 0 seconds.
## Growing trees.. Progress: 40%. Estimated remaining time: 2 minutes, 28 seconds.
## Growing trees.. Progress: 53%. Estimated remaining time: 1 minute, 57 seconds.
## Growing trees.. Progress: 76%. Estimated remaining time: 52 seconds.
```

```
rf_pred <- predict(rf_fit, test_data)$predictions[,2]
```

```
roc.curve(test_data$Class, rf_pred,  
          plotit = TRUE,  
          col = 'green',  
          main = "ROC Curve - Random Forest")
```

```
## Area under the curve (AUC): 0.938
```

```
legend("bottomright", legend = "Random Forest", col = "green", lwd = 2)
```



```
# ----- HYPERPARAMETER TUNING FOR XGBOOST -----
labels <- rose_train_data$Class
y <- ifelse(labels == "Not Fraud", 0, 1)

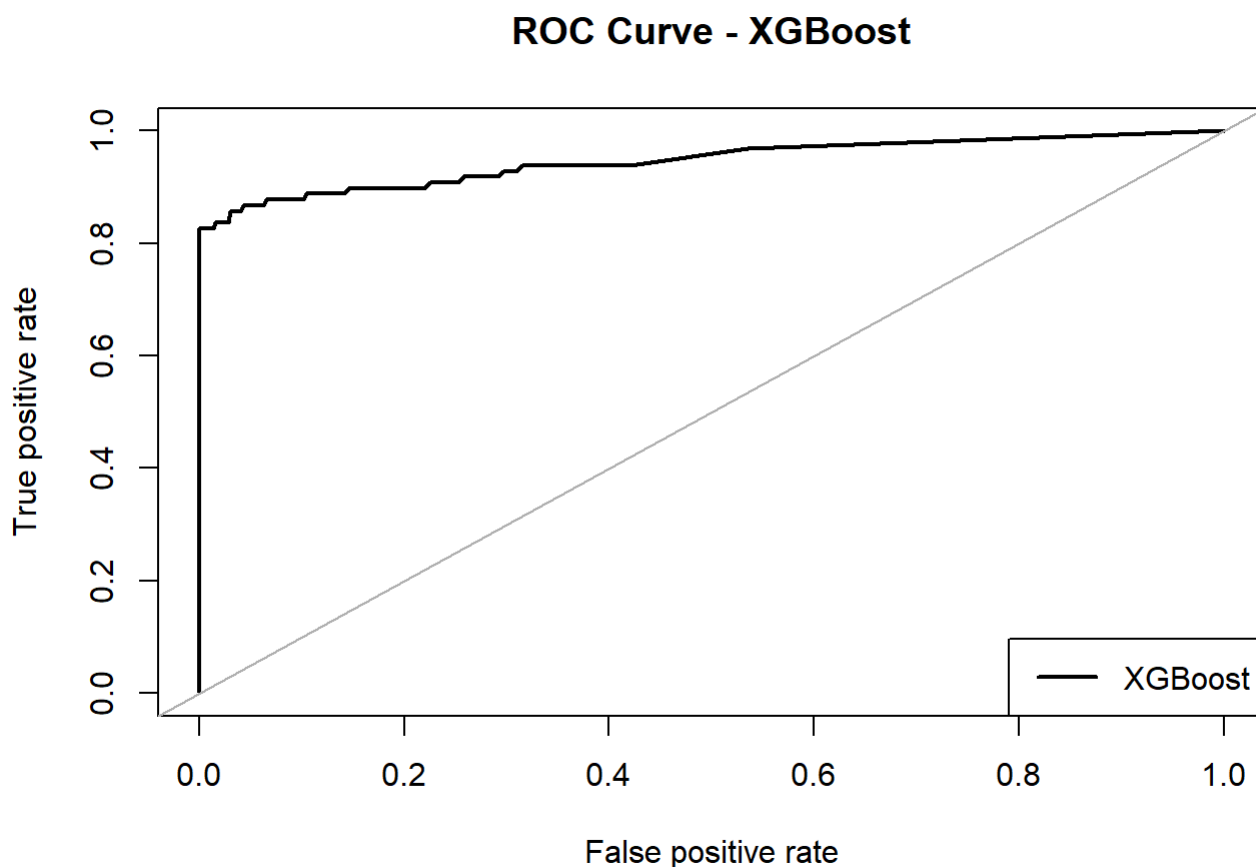
set.seed(42)
xgb <- xgboost(
  data = data.matrix(rose_train_data[,-30]),
  label = y,
  eta = 0.05,
  gamma = 0.1,
  max_depth = 6,
  nrounds = 150,
  objective = "binary:logistic",
  colsample_bytree = 0.7,
  subsample = 0.7,
  verbose = 0,
  nthread = 2
)

xgb_pred <- predict(xgb, data.matrix(test_data[,-30]))

# ROC Curve for XGBoost
roc.curve(test_data$Class, xgb_pred,
  plotit = TRUE,
  col = "black",
  main = "ROC Curve - XGBoost")
```

```
## Area under the curve (AUC): 0.946
```

```
legend("bottomright", legend = "XGBoost", col = "black", lwd = 2)
```



```
# ----- EVALUATION FOR LOGISTIC REGRESSION -----
logistic_pred_class <- ifelse(logistic_predictions > 0.5, "Fraud", "Not Fraud")
logistic_pred_class <- factor(logistic_pred_class, levels = c("Not Fraud", "Fraud"))
logistic_cm <- confusionMatrix(logistic_pred_class, test_data$Class)
logistic_precision <- logistic_cm$byClass['Precision']
logistic_recall <- logistic_cm$byClass['Recall']
logistic_auc <- roc(test_data$Class, logistic_predictions)$auc
```

```
## Setting levels: control = Not Fraud, case = Fraud
```

```
## Setting direction: controls < cases
```

```
# ----- EVALUATION FOR DECISION TREE -----
decisionTree_pred_class <- ifelse(decisionTree_predictions > 0.5, "Fraud", "Not Fraud")
decisionTree_pred_class <- factor(decisionTree_pred_class, levels = c("Not Fraud", "Fraud"))
decisionTree_cm <- confusionMatrix(decisionTree_pred_class, test_data$Class)
decisionTree_precision <- decisionTree_cm$byClass['Precision']
decisionTree_recall <- decisionTree_cm$byClass['Recall']
decisionTree_auc <- roc(test_data$Class, decisionTree_predictions)$auc
```

```
## Setting levels: control = Not Fraud, case = Fraud
## Setting direction: controls < cases
```

```
# ----- EVALUATION FOR RANDOM FOREST -----
rf_pred_class <- ifelse(rf_pred > 0.5, "Fraud", "Not Fraud")
rf_pred_class <- factor(rf_pred_class, levels = c("Not Fraud", "Fraud"))
rf_cm <- confusionMatrix(rf_pred_class, test_data$Class)
rf_precision <- rf_cm$byClass['Precision']
rf_recall <- rf_cm$byClass['Recall']
rf_auc <- roc(test_data$Class, rf_pred)$auc
```

```
## Setting levels: control = Not Fraud, case = Fraud
## Setting direction: controls < cases
```

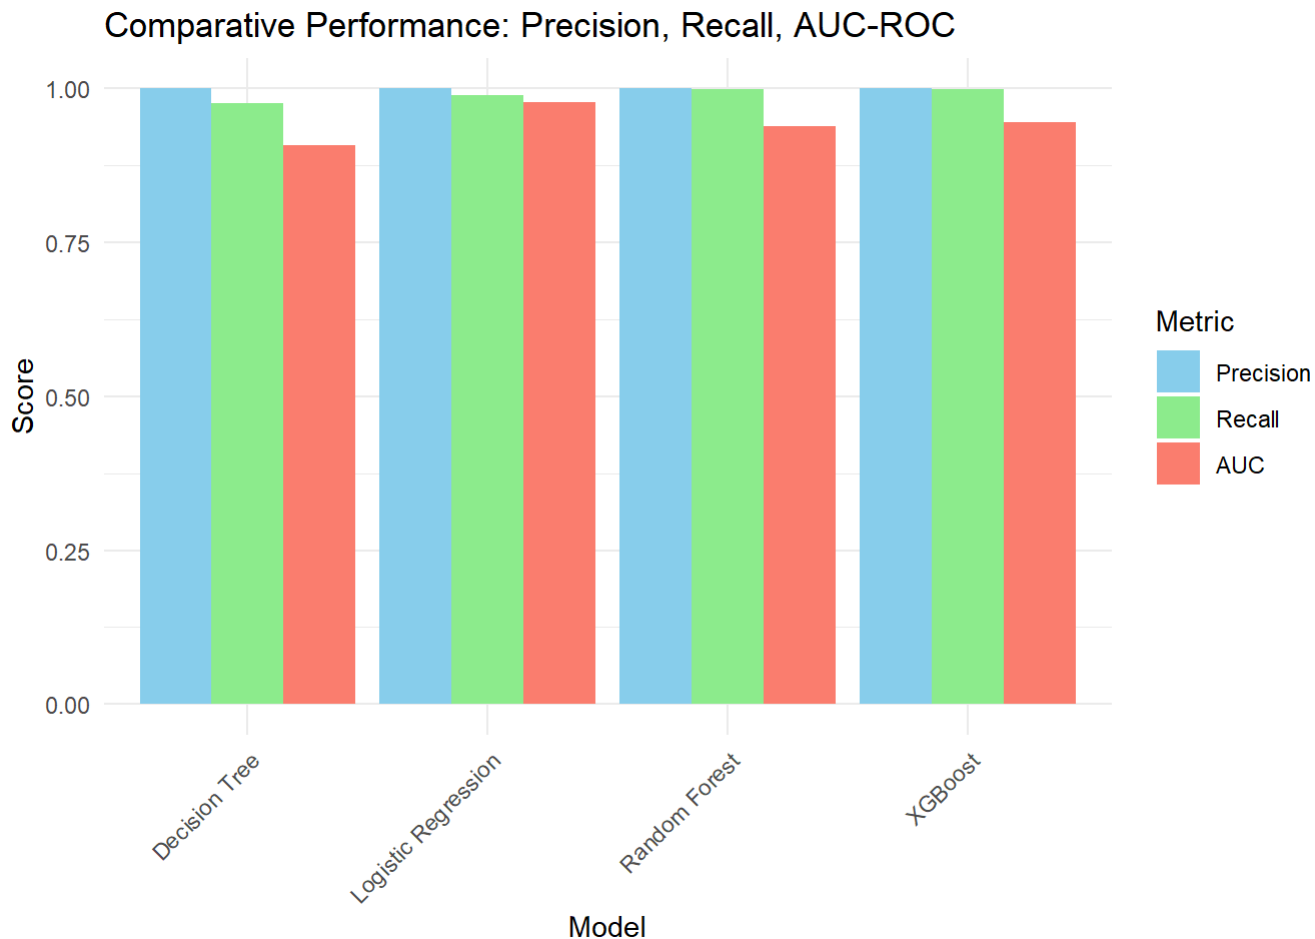
```
# ----- EVALUATION FOR XGBOOST -----
xgb_pred_class <- ifelse(xgb_pred > 0.5, "Fraud", "Not Fraud")
xgb_pred_class <- factor(xgb_pred_class, levels = c("Not Fraud", "Fraud"))
xgb_cm <- confusionMatrix(xgb_pred_class, test_data$Class)
xgb_precision <- xgb_cm$byClass['Precision']
xgb_recall <- xgb_cm$byClass['Recall']
xgb_auc <- roc(test_data$Class, xgb_pred)$auc
```

```
## Setting levels: control = Not Fraud, case = Fraud
## Setting direction: controls < cases
```

```
# ----- COMPARATIVE DATAFRAME -----
model_comparison <- data.frame(
  Model = c("Logistic Regression", "Decision Tree", "Random Forest", "XGBoost"),
  Precision = c(logistic_precision, decisionTree_precision, rf_precision, xgb_precision),
  Recall = c(logistic_recall, decisionTree_recall, rf_recall, xgb_recall),
  AUC = c(logistic_auc, decisionTree_auc, rf_auc, xgb_auc)
)

# Reshape the dataframe for ggplot
model_comparison_long <- reshape2::melt(model_comparison, id.vars = "Model", variable.name = "Metric",
value.name = "Value")

# ----- PLOTTING COMPARATIVE BAR CHART -----
ggplot(model_comparison_long, aes(x = Model, y = Value, fill = Metric)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Comparative Performance: Precision, Recall, AUC-ROC",
    x = "Model",
    y = "Score") +
  scale_fill_manual(values = c("skyblue", "lightgreen", "salmon")) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```




```
# ----- THRESHOLD OPTIMIZATION -----
test_data$Class <- ifelse(test_data$Class == "Fraud", 1, 0)
pred_obj <- prediction(xgb_pred, test_data$Class)
perf <- performance(pred_obj, "tpr", "fpr")

# Find the best threshold (Closest to perfect TPR/FPR balance)
best_threshold_index <- which.max(perf@y.values[[1]] - perf@x.values[[1]])
best_threshold <- perf@alpha.values[[1]][best_threshold_index]

optimized_predictions <- ifelse(xgb_pred > best_threshold, "Fraud", "Not Fraud")
optimized_predictions <- factor(optimized_predictions, levels = c("Not Fraud", "Fraud"))

print(paste("Optimized Threshold:", best_threshold))
```

```
## [1] "Optimized Threshold: 0.0149605302140117"
```

```
# Evaluate Model
test_data$Class <- factor(test_data$Class, levels = c(0, 1), labels = c("Not Fraud", "Fraud"))
conf_matrix <- confusionMatrix(optimized_predictions, test_data$Class)
print(conf_matrix)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  Not Fraud  Fraud
##   Not Fraud    55120    15
##   Fraud        1743     83
##
##              Accuracy : 0.9691
##              95% CI : (0.9677, 0.9705)
##   No Information Rate : 0.9983
##   P-Value [Acc > NIR] : 1
##
##              Kappa : 0.0833
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.96935
##              Specificity : 0.84694
##   Pos Pred Value : 0.99973
##   Neg Pred Value : 0.04545
##   Prevalence : 0.99828
##   Detection Rate : 0.96768
##   Detection Prevalence : 0.96794
##   Balanced Accuracy : 0.90814
##
##   'Positive' Class : Not Fraud
##
```

```

true_labels <- ifelse(test_data$Class == "Fraud", 1, 0)
pred_obj <- prediction(xgb_pred, true_labels)
perf <- performance(pred_obj, "tpr", "fpr")
thresholds <- perf@alpha.values[[1]]
tpr <- perf@y.values[[1]]
fpr <- perf@x.values[[1]]

plot(thresholds, tpr, type = "l", col = "blue", ylim = c(0, 1),
      xlab = "Threshold", ylab = "Rate", main = "TPR and FPR vs. Threshold")
lines(thresholds, fpr, col = "red")
legend("bottomleft", legend = c("TPR", "FPR"), col = c("blue", "red"), lty = 1)

```

